



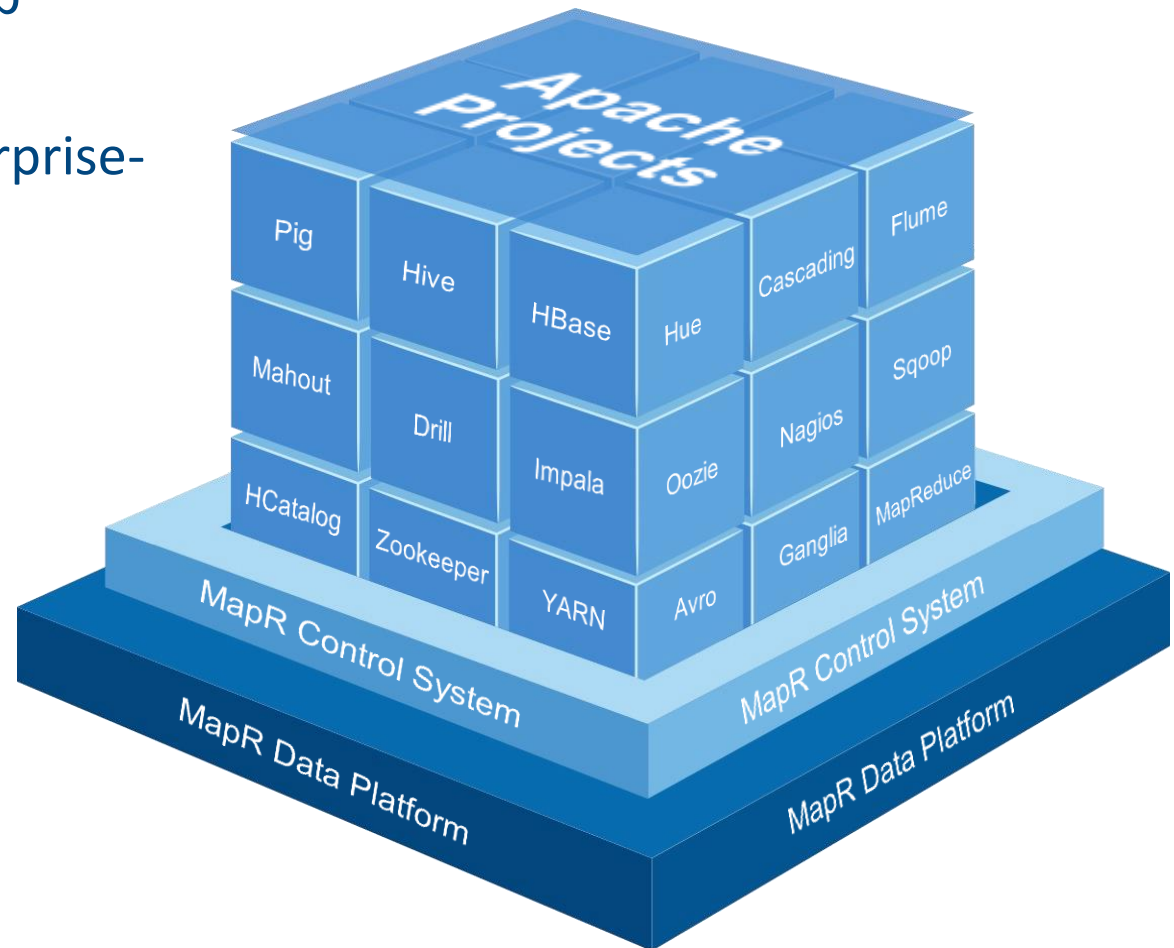
# An Architectural Overview of MapR

**MAPR**<sup>TM</sup>  
TECHNOLOGIES

**M. C. Srivas, CTO/Founder**

# MapR Distribution for Apache Hadoop

- 100% Apache Hadoop
- With significant enterprise-grade enhancements
- Comprehensive management
- Industry-standard interfaces
- Higher performance



# MapR: Lights Out Data Center Ready

## Reliable Compute

- Automated stateful failover
- Automated re-replication
- Self-healing from HW and SW failures
- Load balancing
- Rolling upgrades
- No lost jobs or data
- 99999's of uptime



## Dependable Storage

- Business continuity with snapshots and mirrors
- Recover to a point in time
- End-to-end check summing
- Strong consistency
- Built-in compression
- Mirror between two sites by RTO policy

# MapR does MapReduce (fast)

MAPR<sup>TM</sup>  
TECHNOLOGIES

Google



## TeraSort Record

1 TB in 54 seconds

1003 nodes

## MinuteSort Record

1.5 TB in 59 seconds

2103 nodes

# MapR does MapReduce (faster)

MAPR<sup>TM</sup>  
TECHNOLOGIES

Google



## TeraSort Record

1 TB in 54 seconds

1003 nodes

## MinuteSort Record

1.65 ~~1.5~~ TB in 59 seconds

~~2103~~ nodes

300

# The Cloud Leaders Pick MapR



Amazon EMR is the largest Hadoop provider in revenue and # of clusters

Google chose MapR to provide Hadoop on Google Compute Engine



Deploying OpenStack? MapR partnership with Canonical and Mirantis on OpenStack support.

# How to make a cluster reliable?

1. Make the storage reliable
  - Recover from disk and node failures



# How to make a cluster reliable?

1. Make the storage reliable
  - Recover from disk and node failures
2. Make services reliable
  - Services need to checkpoint their state rapidly
  - Restart failed service, possibly on another node
  - Move check-pointed state to restarted service, using (1) above



# How to make a cluster reliable?

1. Make the storage reliable
  - Recover from disk and node failures
2. Make services reliable
  - Services need to checkpoint their state rapidly
  - Restart failed service, possibly on another node
  - Move check-pointed state to restarted service, using (1) above
3. Do it fast
  - Instant-on ... (1) and (2) must happen very, very fast
  - Without maintenance windows
    - No compactions (eg, Cassandra, Apache HBase)
    - No “anti-entropy” that periodically wipes out the cluster (eg, Cassandra)

# Reliability With Commodity Hardware

- No NVRAM

# Reliability With Commodity Hardware

- No NVRAM
- Cannot assume special connectivity
  - no separate data paths for "online" vs. replica traffic

# Reliability With Commodity Hardware

- No NVRAM
- Cannot assume special connectivity
  - no separate data paths for "online" vs. replica traffic
- Cannot even assume more than 1 drive per node
  - no RAID possible

# Reliability With Commodity Hardware

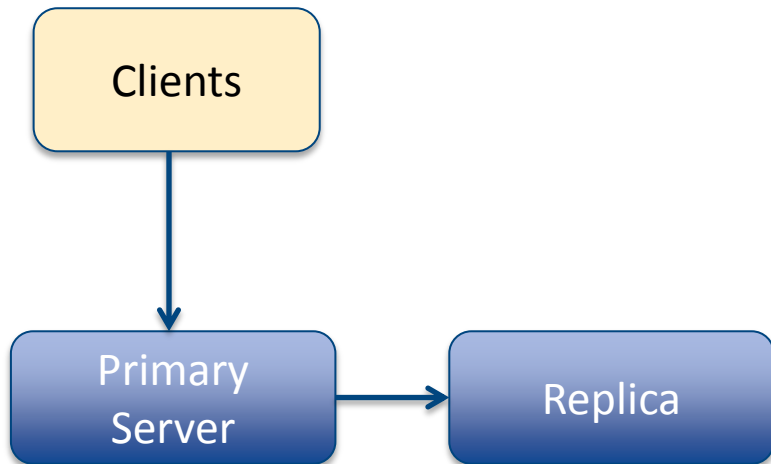
- No NVRAM
- Cannot assume special connectivity
  - no separate data paths for "online" vs. replica traffic
- Cannot even assume more than 1 drive per node
  - no RAID possible
- Use replication, but ...
  - cannot assume peers have equal drive sizes
  - drive on first machine is 10x larger than drive on other?

# Reliability With Commodity Hardware

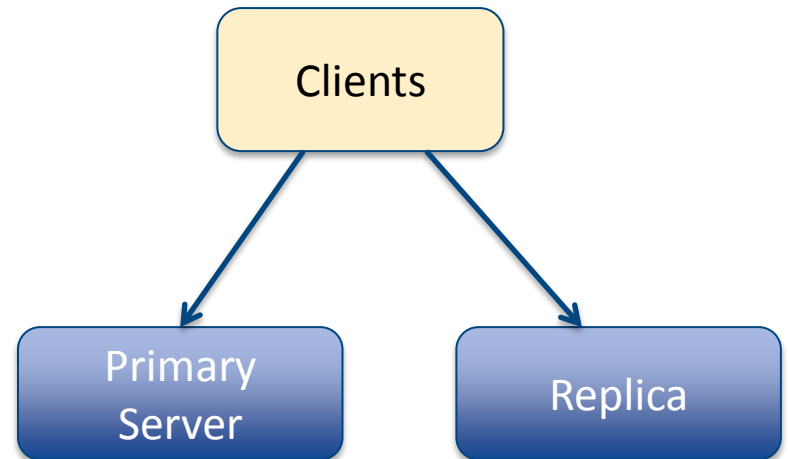
- No NVRAM
- Cannot assume special connectivity
  - no separate data paths for "online" vs. replica traffic
- Cannot even assume more than 1 drive per node
  - no RAID possible
- Use replication, but ...
  - cannot assume peers have equal drive sizes
  - drive on first machine is 10x larger than drive on other?
- No choice but to replicate for reliability

# Reliability via Replication

- Replication is easy, right? All we have to do is send the same bits to the master and replica.



Normal replication, primary forwards

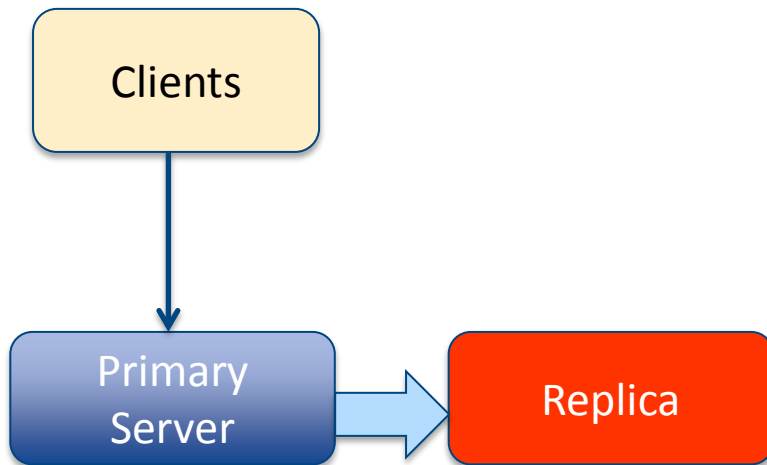


Cassandra-style replication

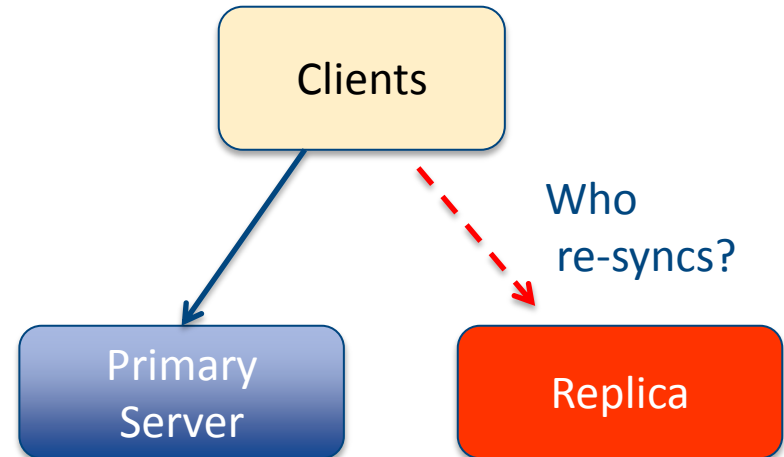


# But crashes occur...

- When the replica comes back, it is stale
  - it must be brought up-to-date
  - until then, exposed to failure



Primary re-syncs replica



Replica remains stale until  
"anti-entropy" process  
kicked off by administrator

# Unless its Apache HDFS ...

- HDFS solves the problem a third way

# Unless its Apache HDFS ...

- HDFS solves the problem a third way
- **Make everything read-only**
  - Nothing to re-sync

# Unless its Apache HDFS ...

- HDFS solves the problem a third way
- **Make everything read-only**
  - Nothing to re-sync
- Single writer, no reads allowed while writing

# Unless its Apache HDFS ...

- HDFS solves the problem a third way
- **Make everything read-only**
  - Nothing to re-sync
- Single writer, no reads allowed while writing
- File close is the transaction that allows readers to see data
  - unclosed files are lost
  - cannot write any further to closed file

# Unless its Apache HDFS ...

- HDFS solves the problem a third way
- **Make everything read-only**
  - Nothing to re-sync
- Single writer, no reads allowed while writing
- File close is the transaction that allows readers to see data
  - unclosed files are lost
  - cannot write any further to closed file
- Real-time not possible with HDFS
  - to make data visible, must close file immediately after writing
  - Too many files is a serious problem with HDFS (a well documented limitation)

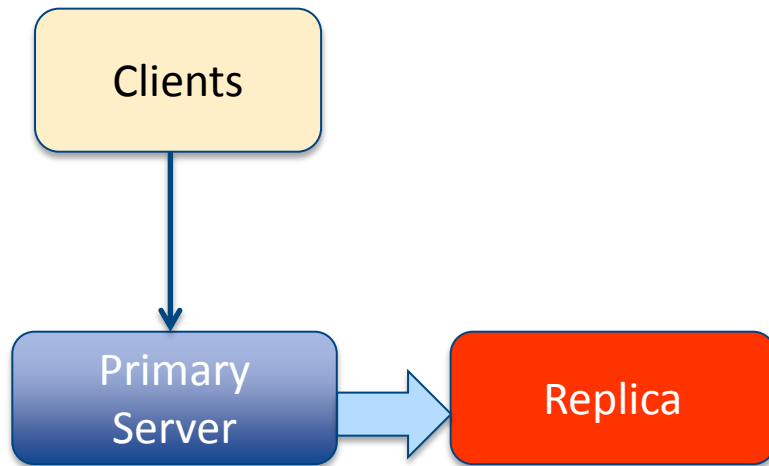
# Unless its Apache HDFS ...

- HDFS solves the problem a third way
- **Make everything read-only**
  - Nothing to re-sync
- Single writer, no reads allowed while writing
- File close is the transaction that allows readers to see data
  - unclosed files are lost
  - cannot write any further to closed file
- Real-time not possible with HDFS
  - to make data visible, must close file immediately after writing
  - Too many files is a serious problem with HDFS (a well documented limitation)
- HDFS therefore cannot do NFS, ever
  - No “close” in NFS ... can lose data any time

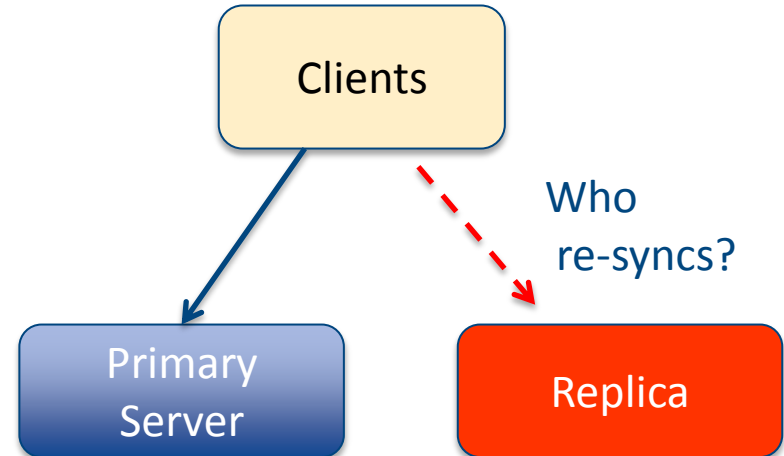


# This is the 21<sup>st</sup> century...

- To support normal apps, need full read/write support
- Let's return to issue: resync the replica when it comes back



Primary re-syncs replica



Replica remains stale until  
"anti-entropy" process  
kicked off by administrator

# How long to re-sync?

- 24 TB / server
  - @ 1000MB/s = 7 hours
  - practical terms, @ 200MB/s = 35 hours

# How long to re-sync?

- 24 TB / server
  - @ 1000MB/s = 7 hours
  - practical terms, @ 200MB/s = 35 hours
- Did you say you want to do this online?

# How long to re-sync?

- 24 TB / server
  - @ 1000MB/s = 7 hours
  - practical terms, @ 200MB/s = 35 hours
- Did you say you want to do this online?
  - throttle re-sync rate to 1/10<sup>th</sup>
  - 350 hours to re-sync (= 15 days)

# How long to re-sync?

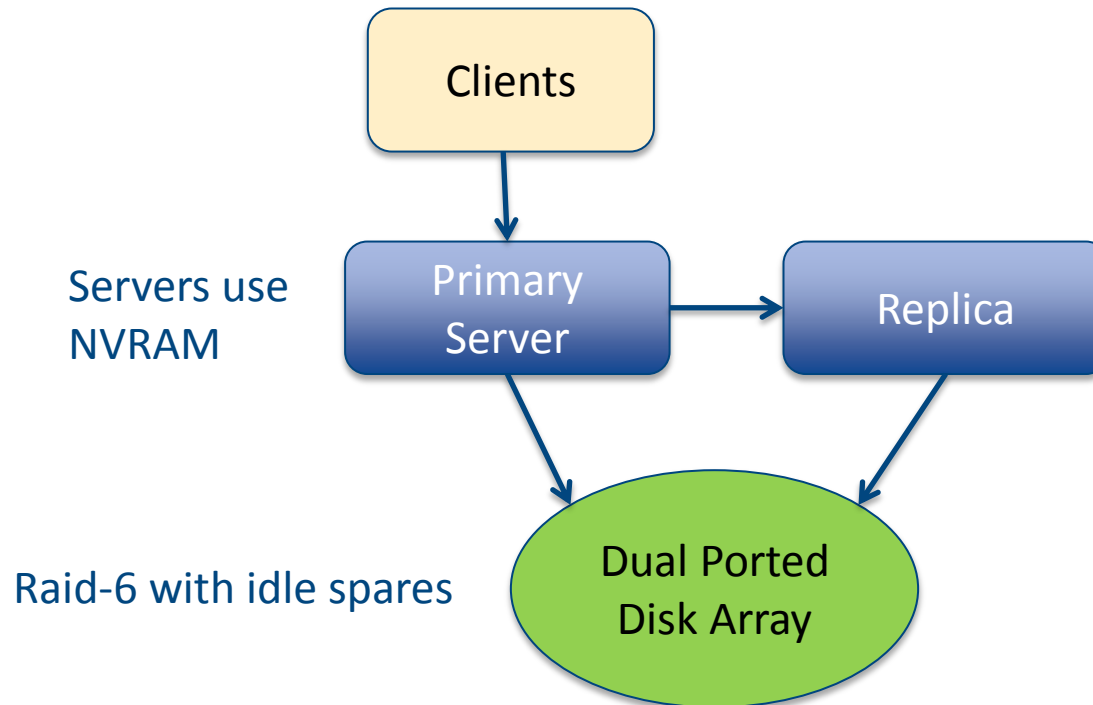
- 24 TB / server
  - @ 1000MB/s = 7 hours
  - practical terms, @ 200MB/s = 35 hours
- Did you say you want to do this online?
  - throttle re-sync rate to 1/10<sup>th</sup>
  - 350 hours to re-sync (= 15 days)
- What is your Mean Time To Data Loss (MTTDL)?

# How long to re-sync?

- 24 TB / server
  - @ 1000MB/s = 7 hours
  - practical terms, @ 200MB/s = 35 hours
- Did you say you want to do this online?
  - throttle re-sync rate to 1/10<sup>th</sup>
  - 350 hours to re-sync (= 15 days)
- What is your Mean Time To Data Loss (MTTDL)?
  - how long before a double disk failure?
  - a triple disk failure?

# Traditional solutions

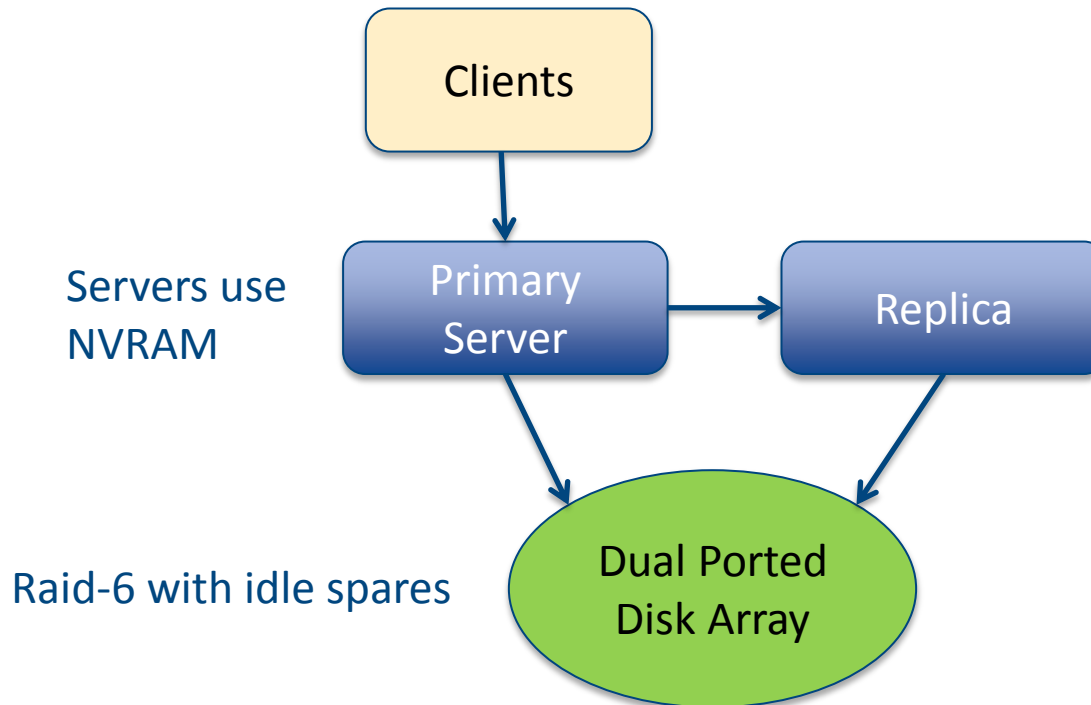
- Use dual-ported disk to side-step this problem





# Traditional solutions

- Use dual-ported disk to side-step this problem




**COMMODITY HARDWARE**

**LARGE SCALE CLUSTERING**

# Traditional solutions

- Use dual-ported disk to side-step this problem




~~COMMODITY HARDWARE~~

~~LARGE SCALE CLUSTERING~~

# Traditional solutions

- Use dual-ported disk to side-step this problem



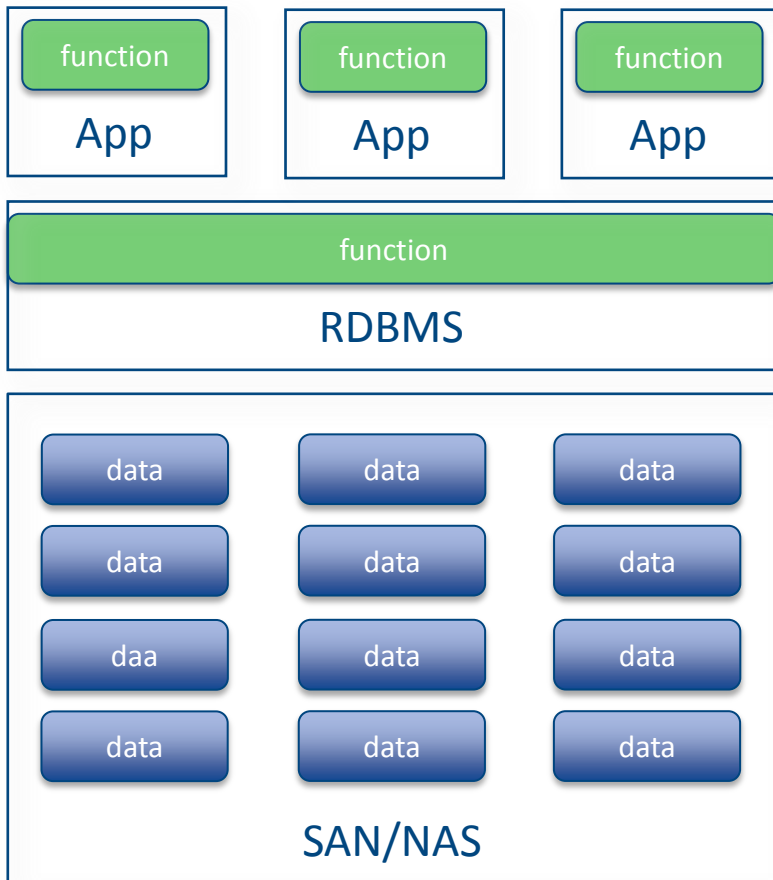
~~COMMODITY HARDWARE~~

~~LARGE SCALE CLUSTERING~~

Large Purchase Contracts, 5-year spare-parts plan

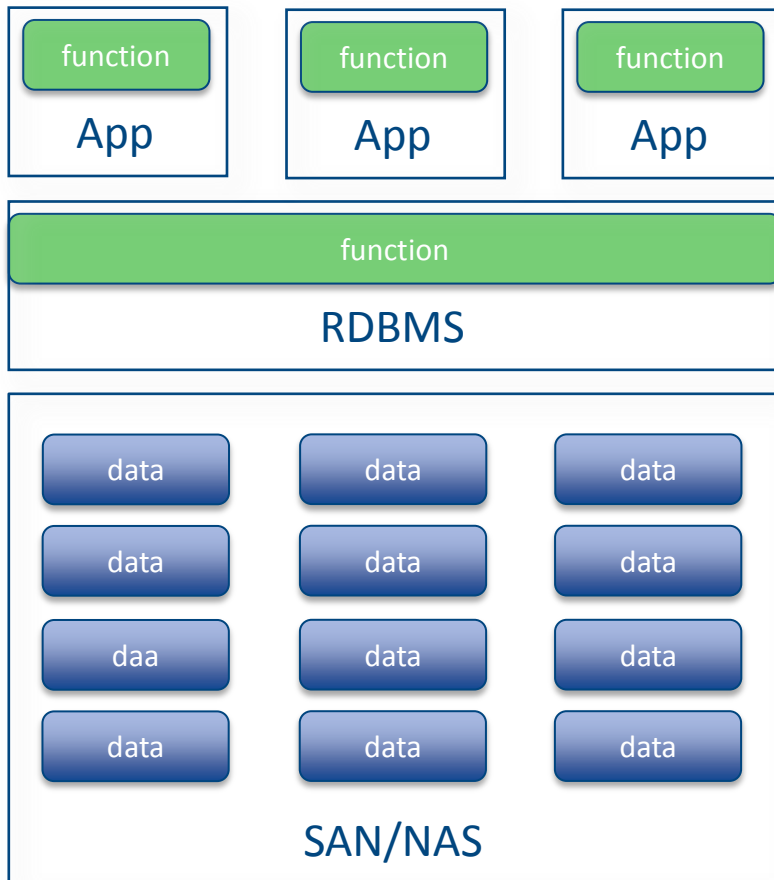
# Forget Performance?

## Traditional Architecture




# Forget Performance?

## Traditional Architecture





## Hadoop




Geographically dispersed also?


# What MapR does

- Chop the data on each node to 1000's of pieces
  - not millions of pieces, only 1000's
  - pieces are called *containers* 




# What MapR does


- Chop the data on each node to 1000's of pieces
  - not millions of pieces, only 1000's
  - pieces are called ***containers*** 







# What MapR does

- Chop the data on each node to 1000's of pieces
  - not millions of pieces, only 1000's
  - pieces are called **containers** 
- Spread replicas of each container across the cluster



# What MapR does


- Chop the data on each node to 1000's of pieces
  - not millions of pieces, only 1000's
  - pieces are called **containers** 
- Spread replicas of each container across the cluster



**Why does it improve things?**


# MapR Replication Example

- 100-node cluster
- each node holds  $1/100^{\text{th}}$  of every node's data




# MapR Replication Example

- 100-node cluster
- each node holds 1/100<sup>th</sup> of every node's data
- when a server dies




# MapR Replication Example

- 100-node cluster
- each node holds 1/100<sup>th</sup> of every node's data
- when a server dies




# MapR Replication Example

- 100-node cluster
- each node holds 1/100<sup>th</sup> of every node's data
- when a server dies



# MapR Replication Example


- 100-node cluster
- each node holds 1/100<sup>th</sup> of every node's data
- when a server dies
- entire cluster resync's the dead node's data






# MapR Replication Example

- 100-node cluster
- each node holds 1/100<sup>th</sup> of every node's data
- when a server dies
- entire cluster resync's the dead node's data




# MapR Re-sync Speed

- 99 nodes re-sync'ing in parallel




# MapR Re-sync Speed

- 99 nodes re-sync'ing in parallel
  - 99x number of drives
  - 99x number of ethernet ports
  - 99x cpu's




# MapR Re-sync Speed

- 99 nodes re-sync'ing in parallel
  - 99x number of drives
  - 99x number of ethernet ports
  - 99x cpu's
- Each is resync'ing 1/100<sup>th</sup> of the data




# MapR Re-sync Speed

- 99 nodes re-sync'ing in parallel
  - 99x number of drives
  - 99x number of ethernet ports
  - 99x cpu's
- Net speed up is about 100x
  - 350 hours vs. 3.5
- Each is resync'ing 1/100<sup>th</sup> of the data



# MapR Re-sync Speed


- 99 nodes re-sync'ing in parallel
  - 99x number of drives
  - 99x number of ethernet ports
  - 99x cpu's
- Net speed up is about 100x
  - 350 hours vs. 3.5
- MTDDL is **100x** better
- Each is resync'ing 1/100<sup>th</sup> of the data



**Why is this so difficult?**

# MapR's Read-write Replication


- Writes are synchronous
- Data is replicated in a "chain" fashion
  - utilizes full-duplex network
- Meta-data is replicated in a "star" manner
  - response time better





# Container Balancing


- Servers keep a bunch of containers "ready to go".
- Writes get distributed around the cluster.



- As data size increases, writes spread more, like dropping a pebble in a pond
- Larger pebbles spread the ripples farther
- Space balanced by moving idle containers


# MapR Container Resync

- MapR is 100% random write  
– very tough problem
- On a complete crash, all replicas diverge from each other
- On recovery, which one should be master?




# MapR Container Resync

- MapR can detect exactly where replicas diverged
  - even at 2000 MB/s update rate
- Resync means
  - roll-back rest to divergence point
  - roll-forward to converge with chosen master
- Done while online
  - with very little impact on normal operations



# MapR does Automatic Resync Throttling

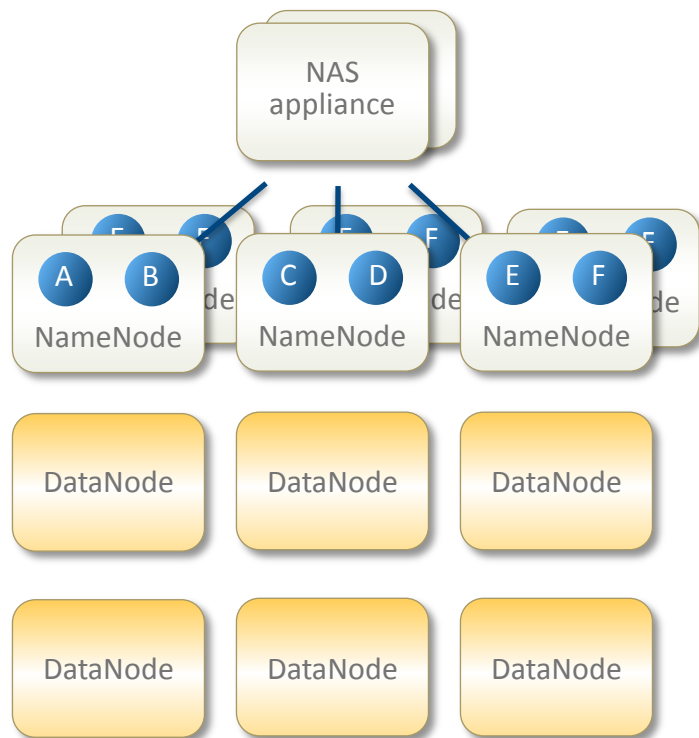
- Resync traffic is “secondary”
- Each node continuously measures RTT to all its peers
- More throttle to slower peers
  - Idle system runs at full speed
- All automatically



**Where/how does MapR exploit this  
unique advantage?**


# MapR's No-NameNode Architecture

## HDFS Federation



- Multiple single points of failure
- Limited to 50-200 million files
- Performance bottleneck
- Commercial NAS required



## MapR (distributed metadata)



- HA w/ automatic failover
- Instant cluster restart
- Up to 1T files (> 5000x advantage)
- 10-20x higher performance
- 100% commodity hardware

# Relative performance and scale

	MapR	Other	Advantage
Rate (creates/s)	14-16K	335-360	40x
Scale (files)	6B	1.3M	4615x

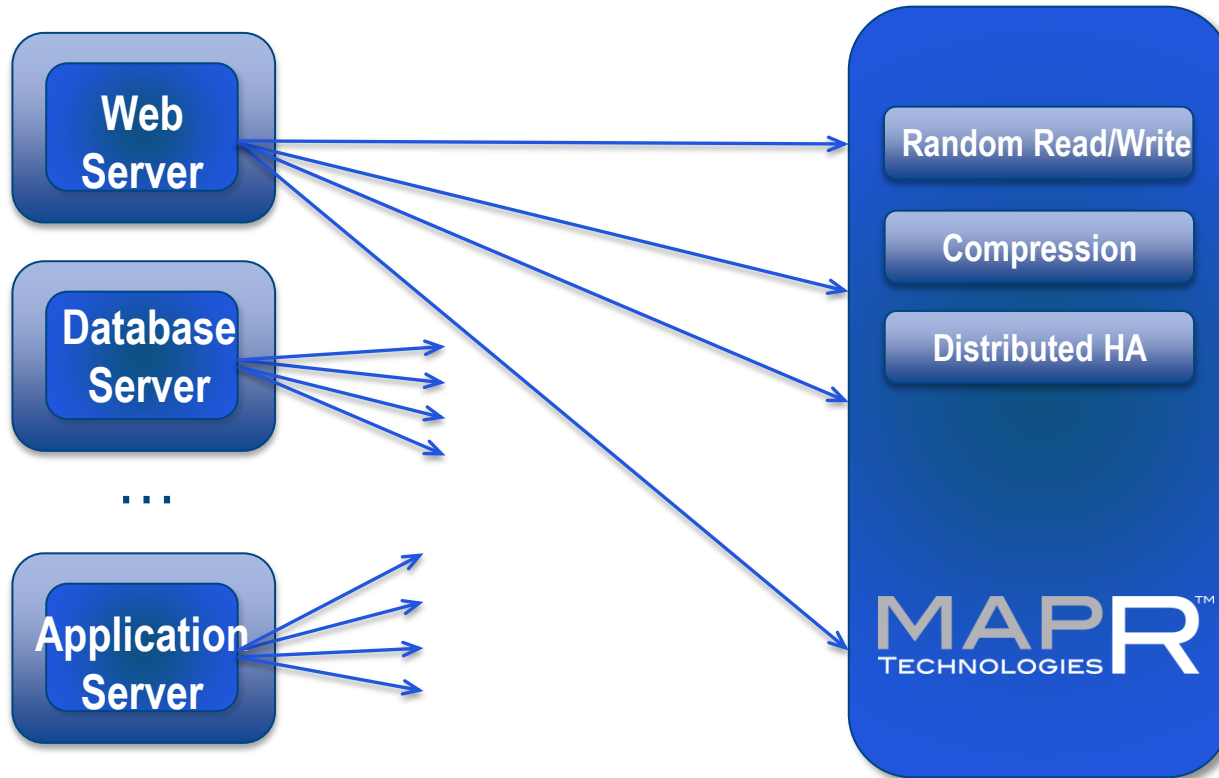


Benchmark: File creates (100B)  
Hardware: 10 nodes, 2 x 4 cores, 24 GB RAM, 12 x 1 TB 7200 RPM

**Where/how does MapR exploit this  
unique advantage?**



# MapR's NFS allows Direct Deposit




Connectors not needed

No extra scripts or clusters to deploy and maintain

**Where/how does MapR exploit this  
unique advantage?**

# MapR Volumes



---

***100K volumes are OK,  
create as many as  
desired!***

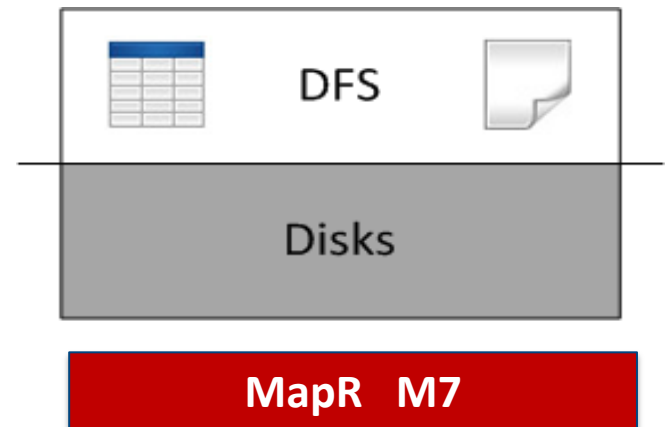
Volumes dramatically simplify the management of Big Data

- Replication factor
- Scheduled mirroring
- Scheduled snapshots
- Data placement control
- User access and tracking
- Administrative permissions

**Where/how does MapR exploit this  
unique advantage?**

# M7 Tables

- M7 tables integrated into storage
  - always available on every node, zero admin
- Unlimited number of tables
  - Apache HBase is typically 10-20 tables (max 100)
- **No compactions**
- **Instant-On**
  - zero recovery time
- 5-10x better perf
- Consistent low latency
  - At 95%-ile and 99%-ile




# M7 vs. CDH: 50-50 Mix (Reads)

**YCSB Mixed (50%Update-50%Read) Test (10Nodes)**

**Source: 2TB (1K RowSize)**

**10-sec Moving Average: Throughput & Read Latency**




# M7 vs. CDH: 50-50 load (read latency)

YCSB Mixed (50%Update-50%Read) Test (10Nodes)

Source: 2TB (1K RowSize)

Read Latency ONLY : 10-sec Moving Average & y-Axis Cap=400msec



**Where/how does MapR exploit this unique advantage?**



# MapR makes Hadoop truly HA

- **ALL** Hadoop components are Highly Available, eg, YARN

# MapR makes Hadoop truly HA

- **ALL** Hadoop components are Highly Available, eg, YARN
- ApplicationMaster (old JT) and TaskTracker record their state in MapR

# MapR makes Hadoop truly HA

- **ALL Hadoop components are Highly Available**, eg, YARN
- ApplicationMaster (old JT) and TaskTracker record their state in MapR
- On node-failure, AM recovers its state from MapR
  - Works even if entire cluster restarted

# MapR makes Hadoop truly HA

- **ALL** Hadoop components are Highly Available, eg, YARN
- ApplicationMaster (old JT) and TaskTracker record their state in MapR
- On node-failure, AM recovers its state from MapR
  - Works even if entire cluster restarted
- All jobs resume **from where they were**
  - Only from MapR

# MapR makes Hadoop truly HA

- **ALL Hadoop components are Highly Available**, eg, YARN
- ApplicationMaster (old JT) and TaskTracker record their state in MapR
- On node-failure, AM recovers its state from MapR
  - Works even if entire cluster restarted
- All jobs resume **from where they were**
  - Only from MapR
- Allows pre-emption
  - MapR can pre-empt any job, without losing its progress
  - ExpressLane™ feature in MapR exploits it

# Where/how can YOU exploit MapR's unique advantage?

- **ALL** your code can easily be scale-out HA

# Where/how can YOU exploit MapR's unique advantage?

- **ALL** your code can easily be scale-out HA
- Save service-state in MapR
- Save data in MapR

# Where/how can YOU exploit MapR's unique advantage?

- **ALL** your code can easily be scale-out HA
- Save service-state in MapR
- Save data in MapR
- Use Zookeeper to notice service failure



# Where/how can YOU exploit MapR's unique advantage?

- **ALL** your code can easily be scale-out HA
- Save service-state in MapR
- Save data in MapR
- Use Zookeeper to notice service failure
- Restart anywhere, data+state will move there automatically

# Where/how can YOU exploit MapR's unique advantage?

- **ALL** your code can easily be scale-out HA
- Save service-state in MapR
- Save data in MapR
- Use Zookeeper to notice service failure
- Restart anywhere, data+state will move there automatically
- That's what we did!

# Where/how can YOU exploit MapR's unique advantage?

- **ALL** your code can easily be scale-out HA
- Save service-state in MapR
- Save data in MapR
- Use Zookeeper to notice service failure
- Restart anywhere, data+state will move there automatically
- That's what we did!
- Only from MapR: HA for Impala, Hive, Oozie, Storm, MySQL, SOLR/Lucene, Kafka, ...

# MapR: Unlimited Scale

# files, # tables	trillions
# rows per table	trillions
# data	1-10 Exabytes
# nodes	10,000+

Build cluster brick by brick, one node at a time

- Use commodity hardware at rock-bottom prices
- Get enterprise-class reliability: instant-restart, snapshots, mirrors, no-single-point-of-failure, ...
- Export via NFS, ODBC, Hadoop and other std protocols